

# Scani5 – End-to-End PRD

**Hybrid Architecture:** Agent in **Go**, Core API in **Go**, Enrichment/Microservices in **Python**.

**Modules Covered:** Agent, Server/API, Discovery/Enrichment, Database, Web App, Security, Deployment, Observability, QA.

**Goal:** Ship a scalable, low-overhead **asset discovery + topology visualization** feature tightly integrated with Scani5 VMDR (vulnerabilities from app list/NVD + external sources), with near real-time updates and enterprise-grade controls.

---

## 0) Executive Summary

Scani5 will extend its VMDR capabilities with a topology subsystem that: (1) augments agents to collect neighbor/service/connection data and full software inventory, (2) ingests and correlates internal telemetry with **OWASP Amass** external discovery, (3) computes and stores a **graph** of assets and interconnections, (4) renders an **interactive visualization** with risk overlays and agent coverage, and (5) continuously enriches nodes with **NVD/Exploit/Advisory** data via Python microservices.

---

## 1) Scope & Objectives

**In scope (v1):** - Scani5 Agent (Go): cross-platform lightweight collectors for **OS, software inventory, services, neighbors, connections, interfaces**. - Internal + external discovery: integrate **Amass** to pull domains/subdomains/IPs/ASNs. - Data platform: PostgreSQL/Timescale for OLTP + time-series; optional Neo4j plugin or graph tables. - Core API (Go): secure ingest, query, export, and WebSocket updates; RBAC/OIDC. - Enrichment services (Python): NVD (CPE/CVE), ExploitDB, vendor advisories; risk scoring. - Web App (React + Cytoscape.js): topology canvas + drilldowns + filters.

**Out of scope (v1):** SNMP/NetFlow, active wide-range port scanning, mobile agents, auto-remediation.

**Success metrics:** - Agent overhead: **≤3% CPU, ≤150 MB RSS** during collection; network payload **≤300 KB** typical. - Freshness: **≥95%** of agent nodes visible within **5 minutes** of check-in. - Scale: render **25k nodes / 200k edges** with <2s pan/zoom interactions (server-side tiling + LOD).

---

## 2) Personas & Top Use Cases

- **SecOps Analyst:** visualize high-risk clusters, see exploitability, spot unagented assets.
- **IT Operations:** validate deployment coverage by site/floor/subnet; plan maintenance.
- **IR/Threat Hunter:** trace lateral movement and blast radius from compromised host.

**Key scenarios**

- 1) Color nodes by **max CVSS**; filter to **site=factory-01 floor=2**; expand neighbors.
- 2) Highlight **unagented** nodes (discovered via neighbors/Amass) to find coverage gaps.
- 3) Click node → details: OS, software list, vuln rollup, services, connections timeline.
- 4) Live rollout: observe agents appearing via **WebSocket delta stream**.
- 5) Export current view to JSON/CSV for audit; API access for SIEM/SOAR.

---

### 3) High-Level Architecture

**Language split - Go:** Agent; Core API (high throughput, mTLS, WS, persistence).

- **Python:** Enrichment microservices (NVD/CPE matching, NLP heuristics, dedupe, scoring).

**Core components** 1. **Scani5 Agent (Go)** → mTLS → **Ingest API (Go)** → **DB (Postgres/Timescale)**.

2. **Amass Connector (Python)** schedules discovery → normalized ingest → DB.

3. **Correlation/Graph Service (Go)** computes edges/clusters/tiles.

4. **Enrichment Workers (Python)** consume app lists → match CPEs → pull CVEs → compute risk → write back.

5. **Web App (React/Cytoscape)** pulls tiles via REST, subscribes to WS for deltas.

**Messaging (optional, at scale):** Kafka/NATS for ingest buffering and enrichment fan-out.

---

### 4) Agent (Go) - Requirements & Design

**Platforms:** Windows 10/11, Server 2016+; Ubuntu 20.04+, RHEL 8+, Amazon Linux 2; macOS 12+.

**Collectors (policy-controlled):** - **System:** hostname, OS name/version, kernel/build, uptime, UUID, site/floor tags. - **Interfaces:** name, MAC, IPs (v4/v6), VLAN, MTU; default route; DNS servers. - **Neighbors:** ARP/NDP, L2 adjacencies; DHCP/Router mac if available. - **Connections (lite):** 5-tuple summaries (src/dst ip/port, proto), counters, first/last seen. - **Services (listening):** TCP/UDP ports → process name/PID; optional binary hash. - **Software inventory:** package/program list + versions (per-OS method below).

- **Resource telemetry:** optional CPU%, mem, payload size (for SRE).

#### OS-specific collection

- **Windows:** WMI/CIM + PowerShell (`Get-CimInstance Win32_Product`) is avoided; use registry & `winget` / `wmic` alternative, `Get-NetTCPConnection`, `Get-NetNeighbor`, `Win32_NetworkAdapter`.
- **Linux:** dpkg/rpm/apk queries; `/proc` for sockets; `ip -j addr`, `ip neigh`, `ss -lntu`.
- **macOS:** `system_profiler SPApplicationsDataType` (parsed), `pkgutil --pkgs`, `arp -an`, `lsof -i -P -n`.

#### Resource controls

- Scheduler with jitter (default 5 min; min 1, max 15).

- Per-collector timeout (2-5s typical).

- Adaptive backoff on high CPU/IO.

- Payload gzip + protobuf/JSON, hard cap (e.g., 1.5 MB).

#### Security & config - mTLS

client auth; cert bootstrap flow; periodic rotation.

- Payload **Ed25519** signature + nonce to prevent replay.

- Remote policy from `/agent/config` (JSON).

- PII guard: hash usernames/process names; redact paths; toggles per policy.

#### Agent Payload (JSON Schema v1)

```
{  
  "$schema": "https://json-schema.org/draft/2020-12/schema",  
  "type": "object",  
  "properties": {  
    "version": {  
      "type": "string",  
      "enum": ["1.0"]  
    },  
    "agent": {  
      "type": "object",  
      "properties": {  
        "name": {  
          "type": "string",  
          "minLength": 1  
        },  
        "version": {  
          "type": "string",  
          "minLength": 1  
        },  
        "os": {  
          "type": "string",  
          "minLength": 1  
        },  
        "arch": {  
          "type": "string",  
          "minLength": 1  
        },  
        "cpu": {  
          "type": "string",  
          "minLength": 1  
        },  
        "memory": {  
          "type": "string",  
          "minLength": 1  
        },  
        "disk": {  
          "type": "string",  
          "minLength": 1  
        },  
        "network": {  
          "type": "string",  
          "minLength": 1  
        },  
        "processes": {  
          "type": "array",  
          "items": {  
            "type": "object",  
            "properties": {  
              "name": {  
                "type": "string",  
                "minLength": 1  
              },  
              "pid": {  
                "type": "number",  
                "minimum": 1  
              },  
              "ppid": {  
                "type": "number",  
                "minimum": 1  
              },  
              "cmdline": {  
                "type": "string",  
                "minLength": 1  
              },  
              "cpu_usage": {  
                "type": "number",  
                "minimum": 0  
              },  
              "mem_usage": {  
                "type": "number",  
                "minimum": 0  
              }  
            }  
          }  
        }  
      }  
    },  
    "config": {  
      "type": "object",  
      "properties": {  
        "interval": {  
          "type": "number",  
          "minimum": 1, "maximum": 15  
        },  
        "timeout": {  
          "type": "number",  
          "minimum": 2, "maximum": 5  
        },  
        "backoff": {  
          "type": "number",  
          "minimum": 1, "maximum": 15  
        },  
        "hard_cap": {  
          "type": "number",  
          "minimum": 1, "maximum": 15  
        },  
        "gzip": {  
          "type": "boolean"  
        },  
        "proto": {  
          "type": "string",  
          "enum": ["protobuf", "json"]  
        },  
        "cert": {  
          "type": "string",  
          "format": "base64"  
        },  
        "key": {  
          "type": "string",  
          "format": "base64"  
        },  
        "policy": {  
          "type": "string",  
          "format": "base64"  
        },  
        "piiguard": {  
          "type": "boolean"  
        }  
      }  
    }  
  }  
}
```

```

"$id": "https://scani5.io/schemas/agent-topology-v1.json",
"title": "Scani5 Agent Topology Payload",
"type": "object",
"required": ["agent_id", "ts", "host", "interfaces"],
"properties": {
    "schema_version": {"type": "string", "const": "1.0"},
    "agent_id": {"type": "string", "format": "uuid"},
    "ts": {"type": "string", "format": "date-time"},
    "host": {
        "type": "object",
        "required": ["hostname", "os", "uuid"],
        "properties": {
            "hostname": {"type": "string"},
            "uuid": {"type": "string"},
            "os": {"type": "object", "properties": {"name": {"type": "string"}, "version": {"type": "string"}, "kernel": {"type": "string"}}},
            "site": {"type": "string"},
            "floor": {"type": "string"}
        }
    },
    "interfaces": {
        "type": "array",
        "items": {"type": "object", "properties": {"name": {"type": "string"}, "mac": {"type": "string"}, "ips": {"type": "array", "items": {"type": "string"}}, "vlan": {"type": "integer"}}},
        "neighbors": {"type": "array", "items": {"type": "object", "properties": {"ip": {"type": "string"}, "mac": {"type": "string"}, "via": {"type": "string"}, "enum": ["arp", "ndp"], "last_seen": {"type": "string", "format": "date-time"}}}},
        "services": {"type": "array", "items": {"type": "object", "properties": {"proto": {"type": "string"}, "enum": ["tcp", "udp"], "port": {"type": "integer"}, "process": {"type": "string"}, "hash": {"type": "string"}, "first_seen": {"type": "string", "format": "date-time"}, "last_seen": {"type": "string", "format": "date-time"}}}},
        "connections": {"type": "array", "items": {"type": "object", "properties": {"src_ip": {"type": "string"}, "src_port": {"type": "integer"}, "dst_ip": {"type": "string"}, "dst_port": {"type": "integer"}, "proto": {"type": "string"}, "count": {"type": "integer"}, "first": {"type": "string", "format": "date-time"}, "last": {"type": "string", "format": "date-time"}}}},
        "apps": {"type": "array", "items": {"type": "object", "required": ["name", "version"], "properties": {"name": {"type": "string"}, "version": {"type": "string"}, "vendor": {"type": "string"}, "install_source": {"type": "string"}, "candidates": {"type": "array", "items": {"type": "string"}}}}},
        "telemetry": {"type": "object", "properties": {"cpu_pct": {"type": "number"}, "mem_mb": {"type": "number"}, "payload_bytes": {"type": "integer"}}},
        "sig": {"type": "string"}
    }
}

```

**Notes:** `apps[*].candidates` may store speculative CPE strings from local heuristics to aid server-side matching.

## 5) Core Server/API (Go) – Requirements & Contracts

- Responsibilities:**
- Authenticate agents via mTLS; validate payload signature/nonce.
  - Normalize & upsert assets, interfaces, services, edges, applications.
  - Serve graph/tiles, summaries, node details; broadcast deltas over WS.
  - Authorize UI/API users via OIDC (Cognito), JWT scopes, RBAC + row-level scoping.

**External Interfaces:** REST + WebSocket; optional gRPC for internal services.

### 5.1 REST/WS Endpoints (OpenAPI snippets)

```
openapi: 3.0.3
info: { title: Scani5 Topology API, version: 1.0.0 }
servers: [{ url: https://api.scani5.internal }]
components:
  securitySchemes:
    mtls: { type: mutualTLS }
    bearerAuth: { type: http, scheme: bearer, bearerFormat: JWT }
  schemas:
    AgentTopologyBatch: { type: object, properties: { items: { type: array, items: { $ref: '#/components/schemas/AgentTopology' }}}} }
    AgentTopology: { $ref: 'https://scani5.io/schemas/agent-topology-v1.json' }
    GraphTile:
      type: object
      properties:
        tile: { type: string }
        nodes: { type: array, items: { $ref: '#/components/schemas/Node' } }
        edges: { type: array, items: { $ref: '#/components/schemas/Edge' } }
    Node:
      type: object
      properties:
        id: { type: string, format: uuid }
        label: { type: string }
        os: { type: string }
        agented: { type: boolean }
        risk: { type: number }
        ips: { type: array, items: { type: string } }
        last_seen: { type: string, format: date-time }
    Edge:
      type: object
      properties:
        src: { type: string, format: uuid }
        dst: { type: string, format: uuid }
        kind: { type: string, enum: [neighbor, connection] }
        weight: { type: integer }
```

```

        last_seen: { type: string, format: date-time }
security: []
paths:
  /ingest/topology:
    post:
      security: [{ mtls: [] }]
      requestBody:
        required: true
        content:
          application/json:
            schema: { $ref: '#/components/schemas/AgentTopologyBatch' }
      responses:
        '202': { description: Accepted }
  /agent/config:
    get:
      security: [{ mtls: [] }]
      parameters: [{ name: agent_id, in: query, schema: { type: string,
format: uuid } }]
      responses: { '200': { description: OK } }
  /topology/summary:
    get:
      security: [{ bearerAuth: [] }]
      parameters:
        - { name: site, in: query, schema: { type: string } }
        - { name: floor, in: query, schema: { type: string } }
      responses: { '200': { description: OK } }
  /topology/graph:
    get:
      security: [{ bearerAuth: [] }]
      parameters:
        - { name: site, in: query, schema: { type: string } }
        - { name: floor, in: query, schema: { type: string } }
        - { name: only_agented, in: query, schema: { type: boolean } }
        - { name: min_severity, in: query, schema: { type: number } }
        - { name: tile, in: query, schema: { type: string } }
      responses:
        '200':
          content:
            application/json:
              schema: { $ref: '#/components/schemas/GraphTile' }
  /topology/node/{id}:
    get:
      security: [{ bearerAuth: [] }]
      parameters: [{ name: id, in: path, required: true, schema: { type:
string, format: uuid } }]
      responses: { '200': { description: OK } }
  /export/topology:
    get:
      security: [{ bearerAuth: [] }]
      responses: { '200': { description: file } }
  /ws/topology:

```

```

get:
  security: [{ bearerAuth: [] }]
  responses: { '101': { description: Switching Protocols } }

```

**Notes:** Keep API response payloads **tile-based** to avoid sending entire graphs.

## 5.2 Ingest & Processing Flow

- 1) Validate mTLS + payload signature; dedupe via nonce cache (Redis).
- 2) Normalize MAC/IP formats; upsert **assets**, **interfaces**, **services**, **apps**.
- 3) Build/merge **edges** (**neighbor** / **connection**) with weights + timestamps.
- 4) Publish **delta events** (Redis pub/sub) → WS broadcaster.
- 5) Queue **apps** for enrichment (Kafka topic: `apps.to_enrich`).

## 5.3 Policy Model (example)

```

{
  "collect": {
    "interval_seconds": 300,
    "neighbors": true,
    "services": true,
    "connections": true,
    "apps": true
  },
  "privacy": { "hash_process_names": true, "redact_paths": true },
  "limits": { "max_neighbors": 5000, "max_connections": 2000 }
}

```

## 6) Enrichment & Discovery (Python microservices)

**Services (FastAPI workers / Celery/Arq):** - **CPE Matcher:** transform `(name, version, vendor)` to candidate **CPEs** using rules + fuzzy match.  
 - **NVD Fetcher:** pull CVEs by CPE from NVD JSON feeds; cache; store **CVSS v3.1**, CWE, references.  
 - **ExploitDB/Vendor Advisories:** scrape/API where allowed; mark **public exploit availability**.  
 - **Risk Scorer:** compute `risk_score` per asset: `max(cvss) + exploit_bonus + age_factor + exposure (services/open ports)`.

### Data contracts

```

{
  "app_key": {"name": "openssl", "version": "1.1.1k", "vendor": "OpenSSL"},
  "cpe_candidates": ["cpe:2.3:a:openssl:openssl:1.1.1k:*:*:*:*:*"],
  "cves": [{"cve": "CVE-2023-12345", "cvss": 8.8, "exploitable": true}],
  "risk": {"score": 9.3, "factors": {"cvss": 8.8, "exploit": 0.3, "exposure": 0.2}}
}

```

**Storage:** tables `apps`, `app_cpe_links`, `cves`, `app_cvss`, `asset_risk` (see DB section).

**Amass Connector** - Scheduled (cron) to run `passive` by default; configurable `active`.

- Parse JSON output; upsert `amass_assets (fqdn, ip, asn, first_seen/last_seen)`; correlate with internal `assets` (by IP/FQDN) → `asset_links (confidence)`.

---

## 7) Database Design (PostgreSQL + Timescale)

```
-- Assets
CREATE TABLE assets (
    id UUID PRIMARY KEY,
    agent_id UUID,
    hostname TEXT,
    os_name TEXT,
    os_version TEXT,
    kernel TEXT,
    site TEXT,
    floor TEXT,
    mac TEXT,
    ips INET[],
    agented BOOLEAN DEFAULT TRUE,
    risk_score NUMERIC(4,1) DEFAULT 0.0,
    last_seen TIMESTAMPTZ NOT NULL
);
CREATE INDEX assets_site_floor_idx ON assets(site, floor);
CREATE INDEX assets_ips_gin ON assets USING GIN (ips);

-- Interfaces
CREATE TABLE interfaces (
    id UUID PRIMARY KEY,
    asset_id UUID REFERENCES assets(id) ON DELETE CASCADE,
    name TEXT, mac TEXT, ips INET[], vlan INT
);

-- Services
CREATE TABLE services (
    id UUID PRIMARY KEY,
    asset_id UUID REFERENCES assets(id) ON DELETE CASCADE,
    proto TEXT, port INT, process TEXT, hash TEXT,
    first_seen TIMESTAMPTZ, last_seen TIMESTAMPTZ
);
CREATE INDEX services_asset_port_idx ON services(asset_id, port);

-- Applications
CREATE TABLE apps (
    id UUID PRIMARY KEY,
    asset_id UUID REFERENCES assets(id) ON DELETE CASCADE,
    name TEXT, version TEXT, vendor TEXT, install_source TEXT,
    first_seen TIMESTAMPTZ, last_seen TIMESTAMPTZ
```

```

);
CREATE INDEX apps_asset_name_idx ON apps(asset_id, name);

-- Observed Edges
CREATE TABLE edges (
    id BIGSERIAL PRIMARY KEY,
    src_asset UUID REFERENCES assets(id) ON DELETE CASCADE,
    dst_asset UUID REFERENCES assets(id) ON DELETE CASCADE,
    kind TEXT CHECK (kind IN ('neighbor', 'connection')),
    weight BIGINT DEFAULT 1,
    first_seen TIMESTAMPTZ,
    last_seen TIMESTAMPTZ
);
CREATE INDEX edges_src_dst_idx ON edges(src_asset, dst_asset, kind);

-- NVD/CPE/CVE
CREATE TABLE cpes (
    id UUID PRIMARY KEY,
    cpe TEXT UNIQUE,
    vendor TEXT, product TEXT, version TEXT
);
CREATE TABLE cves (
    id UUID PRIMARY KEY,
    cve TEXT UNIQUE,
    cvss NUMERIC(3,1), severity TEXT, published TIMESTAMPTZ, last_modified TIMESTAMPTZ,
    has_exploit BOOLEAN DEFAULT FALSE, summary TEXT
);
CREATE TABLE app_cpe_links (
    app_id UUID REFERENCES apps(id) ON DELETE CASCADE,
    cpe_id UUID REFERENCES cpes(id) ON DELETE CASCADE,
    confidence NUMERIC(3,2),
    PRIMARY KEY (app_id, cpe_id)
);
CREATE TABLE app_cves (
    app_id UUID REFERENCES apps(id) ON DELETE CASCADE,
    cve_id UUID REFERENCES cves(id) ON DELETE CASCADE,
    PRIMARY KEY (app_id, cve_id)
);

-- External (Amass)
CREATE TABLE amass_assets (
    id UUID PRIMARY KEY,
    fqdn TEXT,
    ip INET,
    asn INT,
    source TEXT,
    first_seen TIMESTAMPTZ,
    last_seen TIMESTAMPTZ
);
CREATE INDEX amass_fqdn_idx ON amass_assets(fqdn);

```

```

CREATE TABLE asset_links (
    amass_id UUID REFERENCES amass_assets(id) ON DELETE CASCADE,
    asset_id UUID REFERENCES assets(id) ON DELETE CASCADE,
    confidence NUMERIC(3,2),
    PRIMARY KEY (amass_id, asset_id)
);

-- Time-series for agent health
CREATE TABLE agent_checkins (
    agent_id UUID,
    ts TIMESTAMPTZ,
    site TEXT,
    floor TEXT,
    cpu_pct REAL,
    mem_mb REAL,
    payload_bytes INT
);
SELECT create_hypertable('agent_checkins','ts');

```

### Views/rollups

```

CREATE MATERIALIZED VIEW asset_vuln_rollup AS
SELECT a.id AS asset_id, COALESCE(MAX(c.cvss),0) AS max_cvss, COUNT(DISTINCT
c.cve) AS cve_count
FROM assets a
LEFT JOIN apps p ON p.asset_id=a.id
LEFT JOIN app_cves ac ON ac.app_id=p.id
LEFT JOIN cves c ON c.id=ac.cve_id
GROUP BY a.id;

```

## 8) Web App (React + Vite + Cytoscape.js)

**Pages - Topology** (default): infinite canvas w/ mini-map; tiles fetched from [/topology/graph?  
tile=...](/topology/graph?tile=...); WS deltas for live updates.

- **Assets**: searchable table; drill into node details.
- **Risk Explorer**: filters by CVSS/exploitability; heat overlays.
- **Admin**: Policy editor, Amass settings, RBAC, audit logs.

**UI Interactions** - Node click → Drawer with OS, apps, services, max CVSS, last seen, tags.

- Double-click node → expand N neighbors (configurable).
- Lasso select → group/pin; export selection.
- Filters: site, floor, subnet, OS, risk, agented.
- Layouts: COSE (force-directed) default; hierarchical by subnet; geo/floor grid.

**Performance strategies** - Server-side **tile & cluster** aggregation; send only visible subset.

- Client **LOD**: hide labels on zoom-out; edge bundling; cap max render counts.

**Security** - OIDC login (Cognito), JWT in SPA; per-API site scoping; ETag caching.

---

## 9) Security, Privacy, Compliance

- Transport: TLS1.3, **mTLS** for agents, JWT for UI; cert pinning in agent.
  - Payload security: Ed25519 signature + nonce; replay cache in Redis.
  - AuthZ: RBAC (`admin`, `secops`, `itops`, `viewer`); row-level scoping by site.
  - Data protection: encryption at rest; PII minimization; hashing of sensitive fields.
  - Supply chain: signed agent binaries; SBOM; SLSA level targets; CI attestation.
  - Auditing: export events, policy changes, admin actions recorded with actor + timestamp.
- 

## 10) Deployment & DevOps

**Environments:** Dev, Staging, Prod.

**Containerization:** Docker images for Go API and Python workers; SPA built to Nginx.

**Kubernetes (Helm) key objects** - Deployments: `api`, `ws-broadcaster`, `graph-service`,  
`amass-connector`, `cpe-matcher`, `nvd-fetcher`, `risk-scorer`, `web`.  
- Stateful: `postgres` (Timescale), optional `neo4j`.  
- Add-ons: `redis`, optional `kafka`.

**Env/Secrets:** AWS SSM/Secrets Manager; per-env config maps.

**CI/CD:** GitHub Actions/GitLab CI: build, test, SAST, container scan, sign, deploy.

### Reference docker-compose (dev)

```
version: '3.9'
services:
  api:
    image: scani5/api:latest
    ports: ["8080:8080"]
    env_file: .env
    depends_on: [db, redis]
  web:
    image: scani5/web:latest
    ports: ["3000:80"]
    depends_on: [api]
  db:
    image: timescale/timescaledb:pg16
    environment: { POSTGRES_PASSWORD: postgres }
    ports: ["5432:5432"]
  redis:
    image: redis:7
    ports: ["6379:6379"]
  nvd:
    image: scani5/nvd-fetcher:latest
    depends_on: [db]
```

```
risk:  
image: scani5/risk-scorer:latest  
depends_on: [db]
```

## 11) Observability & SLOs

- **Metrics:** ingest TPS, payload sizes, WS subscribers, render tile latency, DB qps.
- **Logs:** structured JSON; correlation IDs; redaction filters.
- **Tracing:** OpenTelemetry across API → DB → enrichment → WS.
- **SLOs:** 99p ingest < 2s; 95p WS fan-out < 500ms; 95p tile fetch < 300ms.

## 12) QA Plan

- **Unit tests:** parsers, dedupe, schema validation, RBAC.
- **Contract tests:** JSON schema for agent payload; OpenAPI for API.
- **Load tests:** 10k agents synthetic; 200k edges; verify latency budgets.
- **Security tests:** mTLS failure, replay attempts, RLS leaks.
- **UI e2e:** Playwright scenarios for topology interactions and filters.

## 13) Implementation Blueprint (Milestones)

### M1 - Agent & Ingest (3-4 weeks) - Go agent collectors (system, interfaces, apps, neighbors, services).

- mTLS bootstrap + signed payloads; `/ingest/topology` upsert path.
- Basic graph build; tiles endpoint; minimal UI prototype.

### M2 - Enrichment & Risk (3-4 weeks) - Python CPE matcher + NVD fetcher; CVE store; risk rollups.

- Node coloring by risk; node detail drawer; exports.

### M3 - Amass & Coverage (2-3 weeks) - Amass connector; external assets; coverage view (unagented highlight).

- Community clustering; performance tuning (LOD/tiles, indexes).

### M4 - Security/RBAC & Admin (2 weeks) - OIDC integration; RBAC; policy editor; audit logs; SLO dashboards.

## 14) Reference Code Stubs

### Go - Agent neighbor collector (sketch)

```
func CollectNeighbors(ctx context.Context) ([]Neighbor, error) {  
    switch runtime.GOOS {  
        case "windows": return collectWinNeighbors()  
        case "darwin": return collectDarwinNeighbors()  
    }
```

```

    default: return collectLinuxNeighbors()
}
}

```

### Go - Ingest handler (gin/fiber/chi)

```

func IngestTopology(w http.ResponseWriter, r *http.Request) {
    // 1) mTLS already enforced at proxy; 2) verify signature+nonce
    var batch AgentTopologyBatch
    if err := json.NewDecoder(r.Body).Decode(&batch); err != nil {
        http.Error(w, "bad", 400)
        return
    }
    for _, item := range batch.Items {
        upsertAsset(tx, item)
        upsertEdges(tx, item)
        enqueueApps(item.Apps)
    }
    w.WriteHeader(http.StatusAccepted)
}

```

### Python - CPE match & NVD fetch (FastAPI worker)

```

from fastapi import FastAPI
app = FastAPI()

@app.post("/enrich/apps")
async def enrich_apps(apps: list[dict]):
    matches = [match_cpe(a) for a in apps]
    cves = fetch_cves_for(matches)
    risk = score(cves, exposure_hint=apps)
    save(matches, cves, risk)
    return {"ok": True}

```

### React - fetch tile

```

const res = await fetch(`/api/topology/graph?site=${site}&tile=${tileId}`, {
    headers: { Authorization: `Bearer ${token}` } });
const tile = await res.json();
cy.add(tile.nodes.map(n => ({ data: n })));
cy.add(tile.edges.map(e => ({ data: e })));

```

## 15) Admin & Governance

- **Config registry** per site/floor (JSON policies); versioned; approvals for changes.
- **Data retention:** connections TTL 30 days; apps/services 180 days; assets forever (soft delete).

- **Exports:** watermark + audit trail; optional PII stripping for reports.
- 

## 16) Risks & Mitigations

- **Volume spikes** → backpressure via queue; batch upserts; payload caps.
  - **CPE ambiguity** → heuristic + human override mapping table.
  - **UI performance** → strict LOD, clustering, pagination/tiles.
  - **Privacy** → hashing, redaction, RBAC scopes; config as code.
  - **Supply chain** → signed builds, SBOM, SLSA provenance, dependency pinning.
- 

## 17) Appendices

**A) Agent Policy JSON Schema**

**B) OpenAPI full spec (expand later)**

**C) Database migrations (Liquibase/Flyway scripts)**

**D) Helm Charts values.yaml templates**

**E) Sample synthetic data generator (Go/Python)**

---

**Deliverable:** This PRD/spec is intentionally **implementation-ready** for AI coding agents: schemas, endpoints, DB DDL, milestones, and code stubs are all provided for scaffolding end-to-end.